

Strategies for real-time event processing

Storm Applied

Sean T. Allen
Matthew Jankowski
Peter Pathirana

FOREWORD BY
Andrew Montalenti





Storm Applied

by Sean T. Allen
Matthew Jankowski
Peter Pathirana

Chapter 5

Copyright 2015 Manning Publications

brief contents

- 1 ■ Introducing Storm 1
- 2 ■ Core Storm concepts 12
- 3 ■ Topology design 33
- 4 ■ Creating robust topologies 76
- 5 ■ Moving from local to remote topologies 102
- 6 ■ Tuning in Storm 130
- 7 ■ Resource contention 161
- 8 ■ Storm internals 187
- 9 ■ Trident 207

5

Moving from local to remote topologies

This chapter covers

- The Storm cluster
- Fault tolerance within a Storm cluster
- Storm cluster installation
- Deploying and running topologies on a Storm cluster
- The Storm UI and the role it plays

Imagine the following scenario. You're tasked with implementing a Storm topology for performing real-time analysis on events logged within your company's system. As a conscientious developer, you've decided to use this book as a guideline for developing the topology. You've built it using the core Storm components covered in chapter 2. You've applied the topology design patterns you learned about in chapter 3 while determining what logic should go into each bolt, and you've followed the steps in chapter 4 to provide at-least-once processing for all tuples coming into your topology. You're ready to hook the topology up to a queue receiving logging events and have it hum along. What do you do next?

You can run your topology locally as in chapters 2, 3, and 4, but doing so won't scale to the data volume and velocity that you're expecting. You need to be able to

deploy your topology to an environment that's built for handling production-level data. This is where the “remote” (also known as “production”) Storm cluster comes into play—an environment built to handle the demands of production-level data.

NOTE As you learned in chapter 1, *volume* refers to the amount of data entering your system and *velocity* refers to the pace at which that data flows through your system.

Running our topologies locally and simulating a Storm cluster within a single process has served our needs so far and is useful for development and testing purposes. But local mode doesn't support the scaling discussed in chapter 3 nor the first-class guaranteed processing we learned about in chapter 4. An actual Storm cluster is needed for both of these.

This chapter will begin by explaining the parts of a Storm cluster and the roles they play, followed by a Q&A session on how Storm provides fault tolerance. We'll then move on to installing a Storm cluster and deploying and running your topologies against the installed cluster. We'll also cover an important tool you can use to make sure your topology is healthy: the Storm UI. Along the way, we'll provide a preview into the tuning and troubleshooting topics that will be covered in chapters 6 and 7.

It all starts with the Storm cluster, so let's expand on our worker nodes discussion from chapter 3.

5.1 The Storm cluster

Chapter 3 scratched the surface of a worker node and how it runs a JVM, which in turn runs executors and tasks. In this section, we're going to go much deeper, starting with the Storm cluster as a whole. A Storm cluster consists of two types of nodes: the master node and the worker nodes. A *master* node runs a daemon called Nimbus, and the *worker* nodes each run a daemon called a Supervisor. Figure 5.1 shows a Storm cluster with one master node and four worker nodes. Storm supports only a single master node, whereas it's likely your cluster will have a different number of worker nodes based on your needs (we'll cover how to determine this number in chapters 6 and 7).

The master node can be thought of as the control center. In addition to the responsibilities listed in figure 5.1, this is where you'd run any of the commands—such as *activate*, *deactivate*, *rebalance*, or *kill*—available in a Storm cluster (more on these commands later in the chapter). The worker nodes are where the logic in the spouts and bolts is executed.

Another big part of a Storm cluster is *Zookeeper*. Storm relies on Apache Zookeeper¹ for coordinating communication between Nimbus and the Supervisors. Any state needed to coordinate between Nimbus and the Supervisors is kept in Zookeeper. As a result, if Nimbus or a Supervisor goes down, once it comes back up it can recover state from Zookeeper, keeping the Storm cluster running as if nothing happened.

¹ <http://zookeeper.apache.org/>




Figure 5.1 Nimbus and Supervisors and their responsibilities inside a Storm cluster

Figure 5.2 shows a cluster of Zookeeper nodes integrated into the Storm cluster. We've removed the worker processes from this figure so you can focus on where Zookeeper fits as it coordinates communication between Nimbus and Supervisors.

Throughout the remainder of the book, any time we mention "Storm cluster," we're referring to the master, worker, and Zookeeper nodes.

Although the master node and Zookeeper are important parts of a Storm cluster, we're going to shift our focus to worker nodes for now. Worker nodes are where the spout and bolt processing occurs, making them the central place for much of our tuning and troubleshooting efforts in chapters 6 and 7.

NOTE Chapters 6 and 7 will explain when you might want to increase the number of worker processes running on a worker node and when and how you might reach a point of diminishing returns. These chapters will also discuss tuning within a worker process, so it makes sense to explain the various parts of a worker process.

5.1.1 The anatomy of a worker node

As mentioned earlier, each worker node has a Supervisor daemon that's tasked with administering the worker processes and keeping them in a running state. If a Supervisor notices that one of the worker processes is down, it will immediately restart it. What's a worker process exactly? We mentioned that it was a JVM, but as you know from chapter 3, there's more to it.




Figure 5.2 The Zookeeper Cluster and its role within a Storm cluster

Each worker process executes a subset of a topology. This means that each worker process belongs to a specific topology and that each topology will be run across one or more worker processes. Normally, these worker processes are run across many machines within the Storm cluster.

In chapter 3, you learned about executors (threads) and tasks (instances of spouts/bolts). We discussed how a worker process (JVM) runs one or more executors (threads), with each thread executing one or more instances of spouts/bolts (tasks). Figure 5.3 illustrates this concept.

Here are the key takeaways:

- A worker process is a JVM.
- An executor is a thread of execution within a JVM.
- A task is an instance of a spout or bolt being run within a thread of execution on the JVM.

Understanding these mappings is extremely important for the purposes of tuning and troubleshooting. For example, chapter 6 answers questions such as why you might want many tasks per executor, so understanding the relationship between an executor and its tasks is essential.

To bring the discussion of a worker node, worker processes, executors, and tasks full circle, let's present them within the context of the credit card authorization topology from chapter 4.

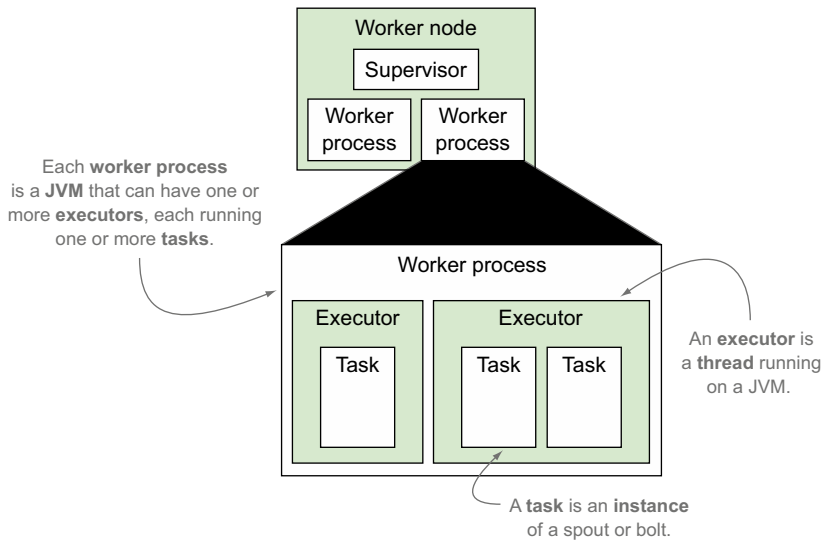


Figure 5.3 A worker process consists of one or more executors, with each executor consisting of one or more tasks.

5.1.2 Presenting a worker node within the context of the credit card authorization topology

In this section we'll present a hypothetical configuration for the credit card authorization topology in order to help you make the connection between the number of worker processes, executors, and tasks in the figures and the code for doing so. This hypothetical configuration can be seen in figure 5.4.

The setup in figure 5.4 would be achieved via the code in the following listing.

Listing 5.1 Configuration for our hypothetical Storm cluster

<p>Set the number of worker processes (JVMs) to 2.</p>	<p>Config config = new Config(); config.setNumWorkers(2); config.setMessageTimeoutSecs(60);</p>	<p>Configure how long each tuple tree has to complete before it gets failed automatically.</p>	<p>Parallelism hint that sets the number of executors (threads) to 1 for this spout, with the default number of tasks (instances) also set to 1</p>
	<p>TopologyBuilder builder = new TopologyBuilder();</p>		
<p>Set the number of executors (threads) to 1.</p>	<p>builder.setSpout("rabbitmq-spout", new RabbitMQSpout(), 1);</p>		
	<p>builder.setBolt("check-status", new VerifyOrderStatus(), 1) .shuffleGrouping("rabbitmq-spout") .setNumTasks(2);</p>		<p>Set the number of tasks (instances) to 2.</p>
<p>Set the number of executors (threads) to 1.</p>	<p>builder.setBolt("authorize-card", new AuthorizeCreditCard(), 1) .shuffleGrouping("check-status") .setNumTasks(2);</p>		<p>Set the number of tasks (instances) to 2.</p>


```
builder.setBolt("notification", new ProcessedOrderNotification(), 1)
    .shuffleGrouping("authorize-card")
    .setNumTasks(1);
```

Set the number of tasks to 1.

Set the number of executors (threads) to 1.

When we set the `numWorkers` in the Config, we're configuring the worker processes desired for running this topology. We don't actually force both worker processes to end up on the same worker node as depicted in figure 5.4. Storm will pick where they end up based on which worker nodes in the cluster have vacant slots for running worker processes.

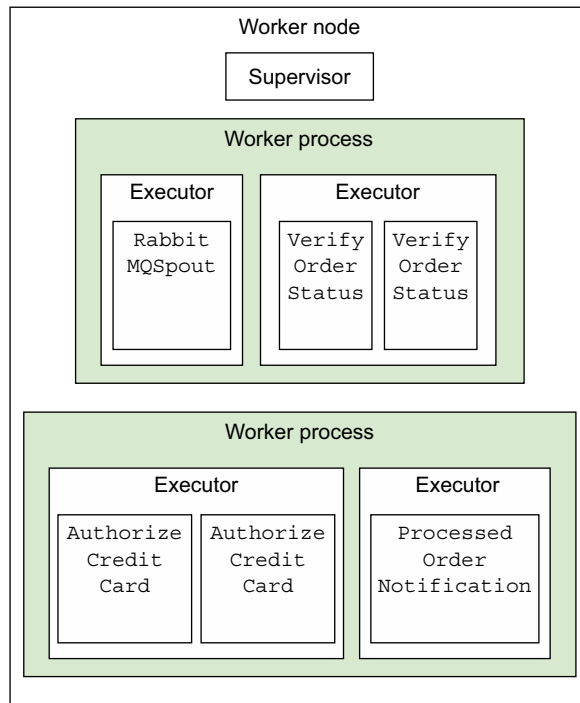


Figure 5.4 A hypothetical breakdown of a worker node with multiple worker processes, executors, and tasks for the credit card authorization topology

Parallelism vs. concurrency: what's the difference?

Parallelism is when two threads are executing simultaneously. *Concurrency* is when at least two threads are making progress on some sort of computation. Concurrency doesn't necessarily mean the two threads are executing simultaneously—something like time slicing may be used to simulate parallelism.

Having revisited the breakdown of a worker node, let's see how Storm provides fault tolerance across the cluster's various parts.

5.2 **Fail-fast philosophy for fault tolerance within a Storm cluster**

Remember the four pieces to the reliability puzzle discussed in chapter 4?

- A reliable data source with a corresponding reliable spout
- An anchored tuple stream
- A topology that acknowledges each tuple as it gets processed or notifies you of failure
- A fault-tolerant Storm cluster infrastructure

We're finally at a point to discuss the last piece, a fault-tolerant Storm cluster infrastructure. The components of a Storm cluster have been designed with fault tolerance in mind. The easiest way to explain how Storm handles fault tolerance is to answer questions in the form of "What does Storm do when x happens?" The most important questions on fault tolerance are addressed in table 5.1.

Table 5.1 Fault tolerance questions and answers

Question	Answer
What if a worker node dies?	Supervisor will restart it and new tasks will be assigned to it. All tuples that weren't fully acked at time of death will be fully replayed by the spout. This is why the spout needs to support replaying (reliable spout) <i>and</i> the data source behind the spout also needs to be reliable (supporting replay).
What if a worker node continuously fails to start up?	Nimbus will reassign tasks to another worker.
What if an actual machine that runs worker nodes dies?	Nimbus will reassign the tasks on that machine to healthy machines.
What if Nimbus dies?	Because Nimbus is being run under supervision (using a tool like daemontools or monit), it should get restarted and continue processing like nothing happened.
What if a Supervisor dies?	Because Supervisors are being run under supervision (using a tool like daemontools or monit), they should get restarted like nothing happened.
Is Nimbus a single point of failure?	Not necessarily. Supervisors and worker nodes will continue to process, but you lose the ability to reassign workers to other machines or deploy new topologies.

You can see that Storm maintains a fail-fast philosophy in the sense that every piece within this infrastructure can be restarted and will recalibrate itself and move on. If tuples were in mid-process during a failure, they'll be failed automatically.

It doesn't matter whether the unit of infrastructure that failed is an instance (task) or a thread (executor) or a JVM (worker process) or a VM (worker node). At each

level, safeguards are in place to ensure that everything gets restarted automatically (because everything runs under supervision).

We've talked about the benefits a Storm cluster provides in terms of parallelism and fault tolerance. How do you go about getting such a cluster up and running?

5.3 Installing a Storm cluster

The Storm wiki does an excellent job describing how to set up a Storm cluster. The steps found on the wiki include these:

- 1 Find information on setting up Zookeeper along with some helpful tips on maintaining the Zookeeper cluster.
- 2 Install the required Storm dependencies on the master and worker machines.
- 3 Download and extract a Storm release to the master and worker machines.
- 4 Configure the master and worker nodes via the `storm.yaml` file.
- 5 Launch the Nimbus and Supervisor daemons under supervision using the Storm script.

We'll cover each of these steps in more detail next.

NOTE What does it mean to run a process under supervision? It means that some supervisory process manages the actual process being run. Therefore, if the process being "supervised" fails, the supervisory process can automatically restart the failed process. This is a key element in providing fault tolerance in Storm.

5.3.1 Setting up a Zookeeper cluster

The steps for setting up a Zookeeper cluster are outside the scope of this book. You can find a thorough explanation of how to install Zookeeper on the Apache Zookeeper project page at <http://zookeeper.apache.org>. Follow those steps to get your cluster up and running.

Keep the following in mind when running Zookeeper:

- Zookeeper is designed to "fail fast," meaning it will shut down if an error occurs that it can't recover from. This isn't desirable within a Storm cluster because Zookeeper coordinates communication between Nimbus and the Supervisors. Because of this, we must have a supervisory process manage the Zookeeper instances so that if a Zookeeper instance does go down, the cluster as a whole can continue handling requests. The supervisory process will handle restarting any failed individual Zookeeper server, allowing the Zookeeper cluster to be self-healing.
- Because Zookeeper is a long-running process, its transaction logs can get quite large. This will eventually result in Zookeeper running out of disk space. Therefore, it's critical to set up some sort of process to compact (and even archive) the data produced in these logs.

5.3.2 Installing the required Storm dependencies to master and worker nodes

The next step is to install the required Storm dependencies to the machines you've dedicated to running Nimbus and the Supervisors. Table 5.2 lists these dependencies.

Table 5.2 External dependencies for Storm master and worker nodes

Dependency	Why it's needed	Link to download
Java 6+	Storm runs on the JVM and the latest version of Storm runs on Java 6.	www.oracle.com/us/technologies/java/overview/index.html
Python 2.6.6	The standard command-line tool for Storm is Python wrapped around Java.	https://www.python.org/downloads/

Once the required external dependencies have been installed to each of the machines hosting Nimbus and the Supervisors, you can install Storm to those machines.

5.3.3 Installing Storm to master and worker nodes

The Storm installations can currently be found at <http://storm.apache.org/downloads.html>. For this book, we used apache-storm-0.9.3. You should download the Storm release zip file to each node and extract the contents of the zip file somewhere on each of the machines. The location is up to you; something like /opt/storm is one example. Figure 5.5 shows the extracted contents in an /opt/storm directory.

There are two files in this figure that we're particularly interested in for this chapter: /opt/storm/bin/storm and /opt/storm/conf/storm.yaml. Let's discuss storm.yaml and its purpose next.

5.3.4 Configuring the master and worker nodes via storm.yaml

The Storm release contains a conf/storm.yaml file that configures the Storm daemons. This file overrides configuration settings found in defaults.yaml.² It's likely that

```

/opt/storm
├── bin
│   ├── storm
│   ├── storm-config.cmd
│   └── storm.cmd
├── conf
│   └── storm.yaml
├── lib
├── logback
├── logs
├── public
├── CHANGELOG.md
├── DISCLAIMER
├── LICENSE
├── NOTICE
├── README.markdown
└── RELEASE

```

Figure 5.5 Extracted contents of a Storm release zip

² You can find defaults.yaml at <https://github.com/apache/storm/blob/master/conf/defaults.yaml>.

you'll want to override at least some of the values; many of the defaults point to "localhost" for a machine name. Table 5.3 lists some of the initial configuration options you may want to override in order to get your Storm cluster up and running.

Table 5.3 storm.yaml properties you may want to override for your Storm installations

Properties	Description	Default value
storm.zookeeper.servers	The lists of hosts in the Zookeeper cluster for your Storm cluster.	storm.zookeeper.servers: - "localhost"
storm.zookeeper.port	Needed if the port your Zookeeper cluster uses is different from the default.	storm.zookeeper.port: 2181
storm.local.dir	The directory that the Nimbus and Supervisor daemons will use for storing small amounts of state. You must create these directories and give them proper permissions on each machine that'll be running Nimbus and workers.	storm.local.dir: "storm-local"
java.library.path	The location of the Java installation.	java.library.path: "/usr/local/lib:/opt/local/lib:/usr/lib"
nimbus.host	The hostname of the Nimbus machine.	nimbus.host: "localhost"
supervisor.slots.ports	For each worker machine, the ports that are used for receiving messages. The number of available ports will determine the number of worker processes Storm runs on each worker machine.	supervisor.slots.ports: - 6700 - 6701 - 6702 - 6703

You'll need to update the configuration on each node in the cluster. Doing so can become tedious if you have a cluster containing several worker nodes. For this reason, we recommend using an external tool such as Puppet³ for automating the deployment and configuration of each node.

5.3.5 Launching Nimbus and Supervisors under supervision

As mentioned earlier, running daemons under supervision is a critical step in setting up a Storm cluster. The supervisory processes allow our system to be fault-tolerant. What does this mean exactly? Why is it needed?

Storm is a fail-fast system, meaning that any Storm process encountering an unexpected error will stop. Storm is designed so that any process can safely stop at any point and recover when the process is restarted. Running these processes under

³ <http://puppetlabs.com/>

supervision allows them to be restarted whenever a failure occurs. Thus, your topologies are unaffected by failures in the Storm daemons. To run the Storm daemons under supervision, execute the following commands:

- *Starting Nimbus*—Run `bin/stormnimbus` under supervision on the master machine.
- *Starting Supervisors*—Run `bin/storm supervisor` under supervision on each worker machine.
- *Storm UI*—Run `bin/storm ui` under supervision on the master machine.

Running the Storm daemons is the last step in setting up a Storm cluster. With everything up and running, your cluster is ready to start accepting topologies. Let's see how you go about getting your topology to run on a Storm cluster.

5.4 **Getting your topology to run on a Storm cluster**

In previous chapters we've run our topologies locally. This approach is fine for learning the fundamentals of Storm. But if one wants to reap the benefits Storm provides (especially along the lines of guaranteed message processing and parallelism), a remote Storm cluster is required. In this section we'll show you how to do this by taking some of the code from the credit card authorization topology in chapter 4 and doing the following:

- Revisit the code for wiring together the topology components
- Show the code for running that topology in local mode
- Show the code for running that topology on a remote Storm cluster
- Show how to package and deploy that code to the remote Storm cluster

5.4.1 **Revisiting how to put together the topology components**

Before we get into the code for running a topology in both local mode and on a remote cluster, let's quickly rehash the code for wiring together the components for the topology from chapter 4, the credit card authorization topology, to provide some context. We've already presented some of this code in section 5.1.2, but the next listing shows it in a more structured format.

Listing 5.2 `CreditCardTopologyBuilder.java` for building the credit card authorization topology

```
public class CreditCardTopologyBuilder {
    public static StormTopology build() {
        TopologyBuilder builder = new TopologyBuilder();

        builder.setSpout("rabbitmq-spout", new RabbitMQSpout(), 1);

        builder.setBolt("check-status", new VerifyOrderStatus(), 1)
            .shuffleGrouping("rabbitmq-spout")
            .setNumTasks(2);

        builder.setBolt("authorize-card", new AuthorizeCreditCard(), 1)
            .shuffleGrouping("check-status")
            .setNumTasks(2);
    }
}
```

```

builder.setBolt("notification", new ProcessedOrderNotification(), 1)
    .shuffleGrouping("authorize-card")
    .setNumTasks(1);

return builder.createTopology();
}
}

```

We've encapsulated the code for building the topology in `CreditCardTopologyBuilder.java` because this code doesn't change, regardless of whether we're running in local mode or on a Storm cluster. This is something we started doing in chapter 3 and the advantage of this approach is it allows us to call the code for building the topology from multiple places without having to duplicate code.

Now that we have the code for building the topology, we'll show you how to take this built topology and run it locally.

5.4.2 Running topologies in local mode

Local mode is useful when developing topologies. It allows you to simulate a Storm cluster in process on your local machine so you can quickly develop and test your topologies. This provides the benefit of a quick turnaround between making a change in code and functionally testing that change in a running topology. There are some drawbacks to local mode, though:

- You can't achieve the parallelism you would with a remote Storm cluster. This makes testing parallelism changes difficult, if not impossible in local mode.
- Local mode won't reveal potential serialization issues when Nimbus attempts to serialize instances of spouts and bolts to the individual worker nodes.

The following listing shows the class, `LocalTopologyRunner`, with a `main()` method that takes the topology we built in listing 5.2 and runs it locally.

Listing 5.3 `LocalTopologyRunner.java`, which runs the topology on a local cluster

```

public class LocalTopologyRunner {
    public static void main(String[] args) {
        StormTopology topology = CreditCardTopologyBuilder.build();

        Config config = new Config();
        config.setDebug(true);

        LocalCluster cluster = new LocalCluster();
        cluster.submitTopology("local-credit-card-topology",
                               config,
                               topology);
    }
}

```

Using the CreditCard-Topology-Builder to build the topology →

← **Normally we run in debug mode when running locally to gain insights into the inner workings of the topology.**

← **Simulates a Storm cluster locally in memory**

← **Submit the topology to the local cluster, passing in the topology name, configuration, and topology.**

The code in this listing should be familiar to you. What we have yet to address is the code for submitting the topology to a remote Storm cluster. Fortunately, this code isn't that much different. Let's take a look.

5.4.3 *Running topologies on a remote Storm cluster*

The code for running your topology remotely is similar to running it locally. The only difference is the code for submitting the topology to a cluster. You may also want a slightly different configuration as well, because local mode doesn't support some of the things (parallelism and guaranteed message processing) that a remote cluster supports. The next listing shows this code in a class we call `RemoteTopologyRunner`.

Listing 5.4 `RemoteTopologyRunner`, which submits the topology to a remote cluster

```

public class RemoteTopologyRunner {
    public static void main(String[] args) {
        StormTopology topology = CreditCardTopologyBuilder.build();

        Config config = new Config();
        config.setNumWorkers(2);
        config.setMessageTimeoutSecs(60);

        StormSubmitter.submitTopology("credit-card-topology",
                                     config,
                                     topology);
    }
}

```

Using the CreditCard-Topology-Builder to build the topology

Set the number of worker processes (JVMs) to 2. This is a config item that we usually tweak only when running a topology on a remote cluster.

Configure how long each tuple tree has to complete before it gets failed automatically.

Using StormSubmitter to submit the topology to the remote cluster, passing in topology name, configuration, and the topology

You'll notice that the only differences are a slightly different configuration and using `StormSubmitter.submitTopology` instead of `LocalCluster.submitTopology`.

NOTE We've wrapped the building, running locally, and running remotely of our topology within three different classes (`CreditCardTopologyBuilder`, `LocalTopologyRunner`, and `RemoteTopologyRunner`). Although you can set up your code however you please, we've found this split works well and we use it across all our topologies.

Now that we've written the code for running our topologies on a remote cluster, let's shift our attention to physically getting that code onto the Storm cluster so it can be run.

5.4.4 *Deploying a topology to a remote Storm cluster*

What does it mean to "deploy" your topology to a Storm cluster? By "deploy," we mean physically copying a JAR containing your topology's compiled code to the cluster so it

can be run. You will need to deploy your topology from a machine with a properly-configured Storm installation. Figure 5.6 provides a refresher on the extracted contents of a Storm release zip.

You'll want to make sure you update the `/opt/storm/conf/storm.yaml` file so the `nimbus.host` property is set to the proper location. We're also interested in the `/opt/storm/bin/storm` file for this step: this is the executable you'll run in order to deploy your topology JAR to the remote cluster. Figure 5.7 shows the command you'd run to deploy your topology. You'll notice in the figure that we reference the full location for the storm executable via `/opt/storm/bin/storm`. If you don't want to do this, put `/opt/storm/bin` on your `PATH` and you can directly reference the `storm` command from anywhere on your machine.

After you execute the command in figure 5.7, your topology will be up and running on the Storm cluster. Once your topology is running, how do you know it's actually working and processing data as expected? This is where you'd look to the Storm UI, which is discussed next.

```

/opt/storm
├── bin
│   ├── storm
│   ├── storm-config.cmd
│   ├── storm-local
│   └── storm.cmd
├── conf
│   └── storm.yaml
├── lib
├── logback
├── logs
├── public
├── CHANGELOG.md
├── DISCLAIMER
├── LICENSE
├── NOTICE
├── README.markdown
└── RELEASE

```

Figure 5.6 Extracted contents of a Storm release zip

```

/opt/storm/bin/storm jar <path-to-topology-jar> <topology-main-class>

```

The command that takes care of connecting to Nimbus and uploading the topology JAR file.

The physical location of the topology JAR file.

Fully qualified name of the main class where the topology is being submitted to the cluster. In our example this would be the fully qualified name of `RemoteTopologyRunner.java`.

Figure 5.7 The command for deploying your topology to a Storm cluster

5.5 The Storm UI and its role in the Storm cluster

The Storm UI is the central place to find diagnostics on the Storm cluster and individual topologies. As mentioned in section 5.3.5, running the command `/bin/storm ui` on Nimbus will start the Storm UI. Two properties in `defaults.yaml` affect where to find the Storm UI:

- 1 `nimbus.host`—The hostname of the Nimbus machine
- 2 `ui.port`—The port number to serve up the Storm UI (defaults to 8080)

Once it's running, enter `http://{nimbus.host}:{ui.port}` in a web browser to get to the Storm UI.

The Storm UI has several sections:

- The Cluster Summary screen
- The individual topology summary screen
- Screens for each of the spouts and bolts

Each screen shows information related to different parts of the Storm cluster at varying levels of granularity. The Cluster Summary screen is related to the Storm cluster as a whole, as seen in figure 5.8.

Clicking on a particular topology link (such as `github-commit-count` in figure 5.8) takes you to a topology summary screen. This screen shows information related to the specific topology, as you can see in figure 5.9.

Let's delve into each screen in more detail next.

5.5.1 Storm UI: the Storm cluster summary

The Storm Cluster Summary consists of four parts, as shown in figure 5.10.

Each section of this screen is explained in more detail in the following sections.

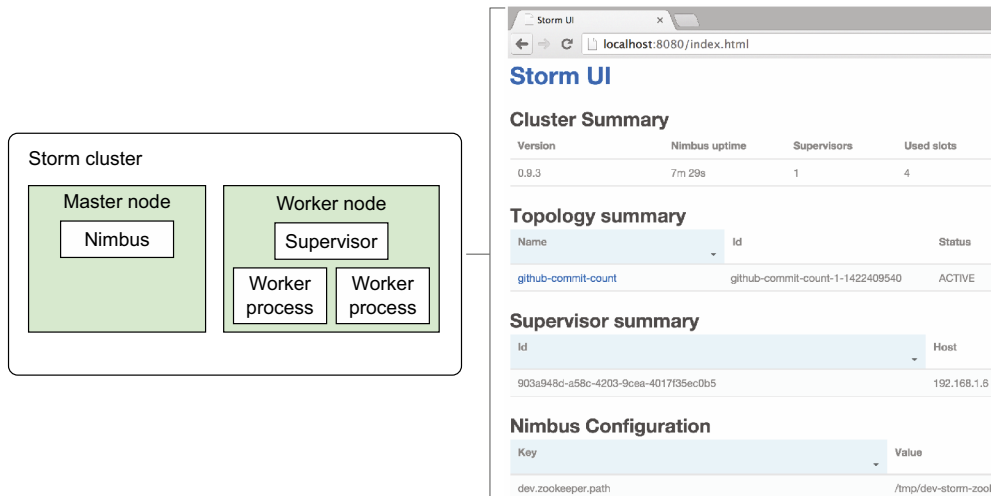


Figure 5.8 The Cluster Summary screen shows details for the entire Storm cluster.

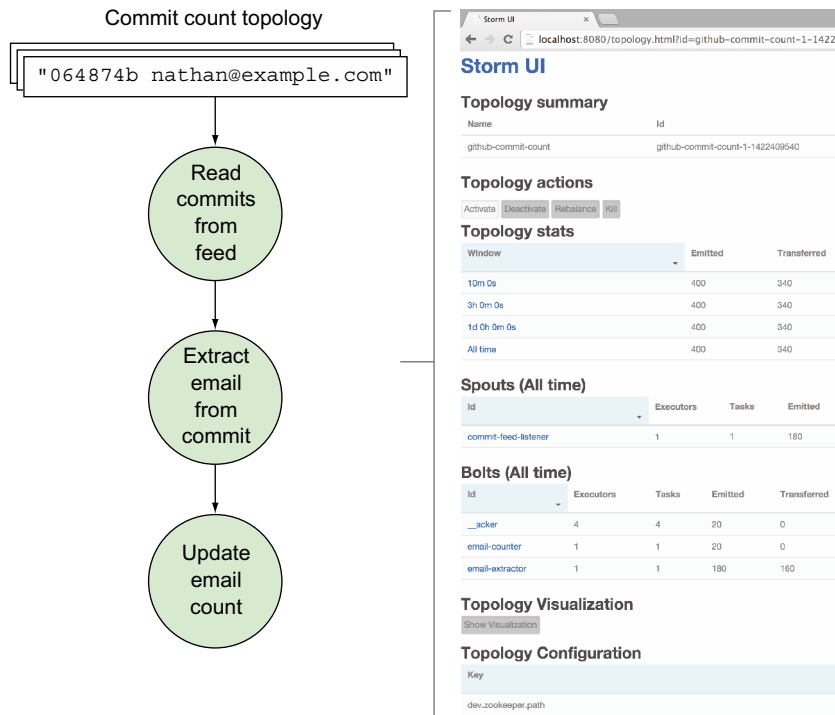


Figure 5.9 The Topology summary screen shows details for a specific topology.

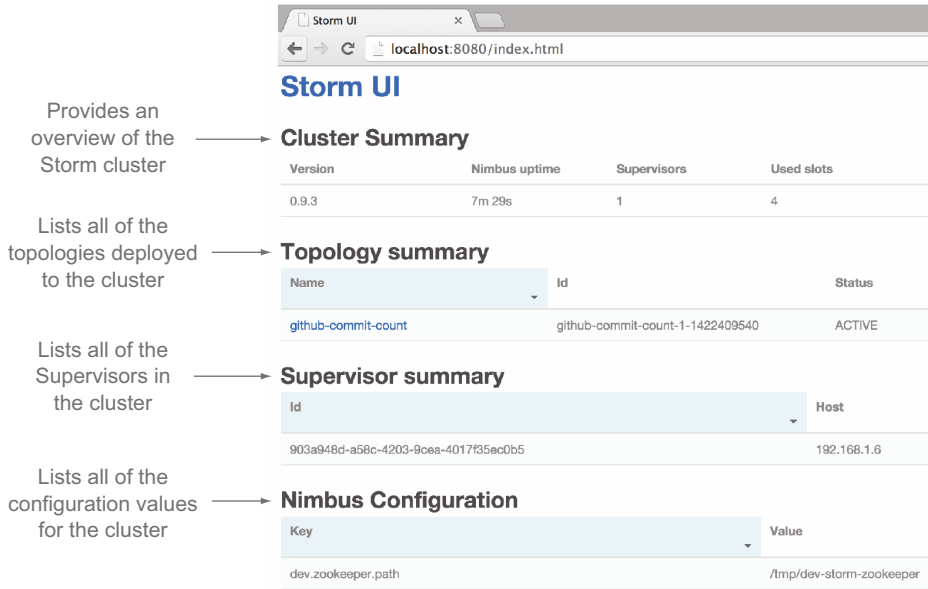


Figure 5.10 The Cluster Summary screen of the Storm UI

CLUSTER SUMMARY

The Cluster Summary section provides a small but useful overview of your Storm cluster. You'll notice the term *slots* in figure 5.11. A slot corresponds to a worker process, so a cluster with two used slots means there are two worker processes running on that cluster. Figure 5.11 provides more detail on each of the columns in this section.

TOPOLOGY SUMMARY

The Topology summary lists all of the topologies deployed to the cluster. Figure 5.12 provides more detail on what information you see in this section.

Cluster Summary

Version	Nimbus uptime	Supervisors	Used slots	Free slots	Total slots	Executors	Tasks
0.9.3	23m 39s	1	2	2	4	9	9

The version of Storm the cluster is running.

The length of time Nimbus has been running.

The number of Supervisor nodes in the cluster.

Number of used slots. This means two worker processes are currently running in the cluster.

Number of free slots. This means that two more worker processes can be run in the cluster.

Number of total slots. This means the cluster can run a total of four worker processes.

The number of executors (threads) being used across the cluster.

Number of tasks (spout/bolt instances) being used across the cluster.

Figure 5.11 The Cluster Summary section on the Cluster Summary screen of the Storm UI

Topology summary

Name	Id	Status	Uptime	Num workers	Num executors	Num tasks
credit-card-authorization	credit-card-authorization-3-1401041926	ACTIVE	18s	1	5	5
github-commit-count	github-commit-count-1-1401041352	ACTIVE	9m 52s	1	4	4

The names of the topologies, defined in the StormSubmitter.submitTopology method.

The IDs assigned to the topologies by Storm.

The current status of the topologies.

The length of time the topologies have been running.

The number of worker processes (JVMs) for the topologies.

The number of executors (threads) for the topologies.

The number of tasks (spout/bolt instances) for the topologies.

Figure 5.12 The Topology summary section on the Cluster Summary screen of the Storm UI

Supervisor summary

Id	Host	Uptime	Slots	Used slots
459bb8f0-92cf-48b1-bdfc-cd99f0c5a3e8	192.168.1.4	23m 9s	4	2

The ID assigned to the Supervisor by Storm.

The length of time the Supervisor node has been running.

The IP address of the Supervisor node.

Number of used slots for the Supervisor. Here two worker processes are running on the Supervisor.

Number of slots for the Supervisor. This Supervisor can run a total of four worker processes.

Figure 5.13 The Supervisor summary section on the Cluster Summary screen of the Storm UI

SUPERVISOR SUMMARY

The Supervisor summary lists all the Supervisors in the cluster. Again, you'll notice the term slots in figure 5.13. This corresponds to a worker process on a particular Supervisor node. Figure 5.13 provides more detail on what information you see in this section.

NIMBUS CONFIGURATION

The Nimbus Configuration lists the configuration defined in defaults.yaml and any overridden values in storm.yaml. Figure 5.14 provides more detail on what information you see in this section.

Nimbus Configuration

Key	Value
dev.zookeeper.path	/tmp/dev-storm-zookeeper
drpc.childopts	-Xmx768m
drpc.invocations.port	3773
drpc.port	3772
drpc.queue.size	128
drpc.request.timeout.secs	600
drpc.worker.threads	64

The Nimbus Configuration item.

The value of the Nimbus Configuration item defined in defaults.yaml or overridden in storm.yaml.

Figure 5.14 The Nimbus Configuration section on the Cluster Summary screen of the Storm UI

Having covered the Cluster Summary screen, let's dive into what the screen for an individual topology looks like. You can access this screen by clicking on a given topology name in the list of topologies.

5.5.2 Storm UI: individual Topology summary

The sections of the individual Topology summary screen can be seen in figure 5.15.

Each section of this screen is explained in more detail in the following sections.

Shows an overview of the topology → **Topology summary**

Name	Id
github-commit-count	github-commit-count-1-1422409540

Provides a UI for performing several Storm commands → **Topology actions**

Shows general statistics at the topology level → **Topology stats**

Window	Emitted	Transferred
10m 0s	400	340
3h 0m 0s	400	340
1d 0h 0m 0s	400	340
All time	400	340

Shows statistics for all spouts in the topology → **Spouts (All time)**

Id	Executors	Tasks
commit-feed-listener	1	1

Shows statistics for all bolts in the topology → **Bolts (All time)**

Id	Executors	Tasks	Emitted	Transferred
_acker	4	4	20	0
email-counter	1	1	20	0
email-extractor	1	1	180	160

Lists all of the configuration values for the topology → **Topology Configuration**

Key
dev.zookeeper.path

Figure 5.15 The Topology summary screen of the Storm UI

Topology summary

Name	Id	Status	Uptime	Num workers	Num executors	Num tasks
github-commit-count	github-commit-count-1-1401041352	ACTIVE	12m 34s	1	4	4

The name of the topology, defined in the StormSubmitter.submitTopology method

The ID assigned to the topology by Storm

The current status of the topology

The length of time the topology has been running

The number of worker processes (JVMs) for the topology

The number of executors (threads) for the topology

The number of tasks (spout/bolt instances) for the topology

Figure 5.16 The Topology summary section on the Topology summary screen of the Storm UI

TOPOLOGY SUMMARY

The Topology summary provides a small but useful overview of the topology being observed. Figure 5.16 provides more detail on each of the individual columns in this section.

TOPOLOGY ACTIONS

The Topology actions section provides a UI for activating, deactivating, rebalancing, and killing your topology. Figure 5.17 describes these actions in more detail.

TOPOLOGY STATS

The Topology stats section provides some general statistics at the topology level. These statistics can be shown for all time, the past 10 minutes, the past 3 hours, or the past day. The time interval that's selected is also applied to the spout and bolt stats sections, which are described next. Figure 5.18 provides more detail on the information you see in this section.

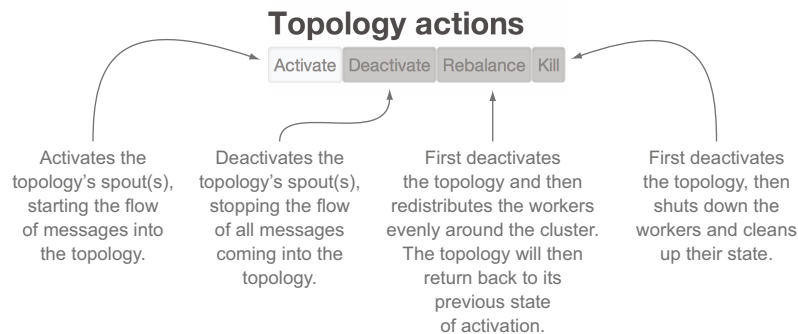


Figure 5.17 The Topology actions section on the Topology summary screen of the Storm UI

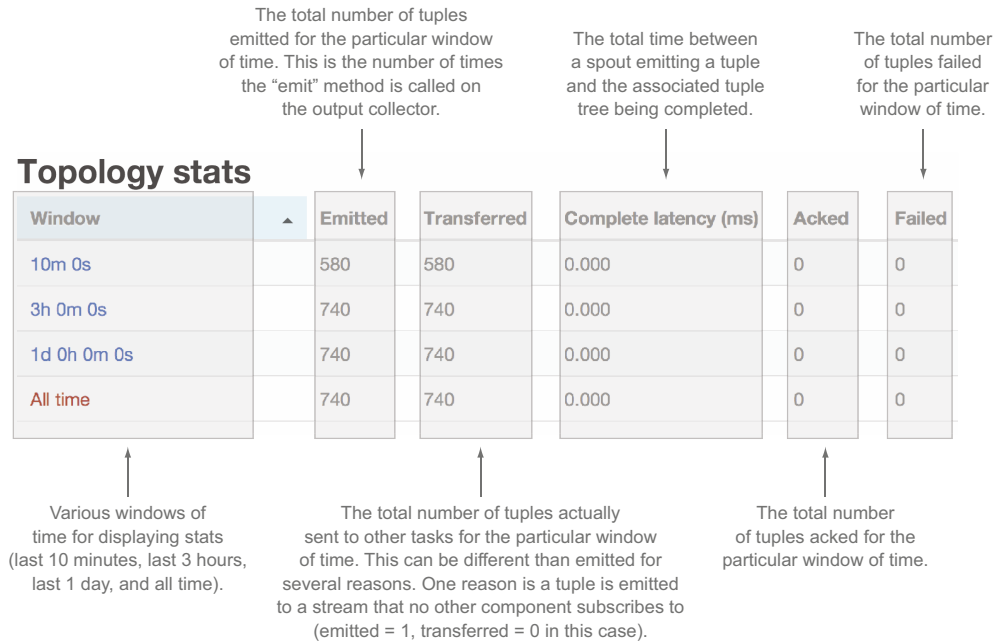


Figure 5.18 The Topology stats section on the Topology summary screen of the Storm UI

SPOUT STATS

The Spouts section shows the statistics for all spouts in the topology. The statistics are for the window of time that's selected in the Topology Stats section (all time, the past 10 minutes, the past 3 hours, or the past day). Figure 5.19 provides more detail on the information you see in this section.

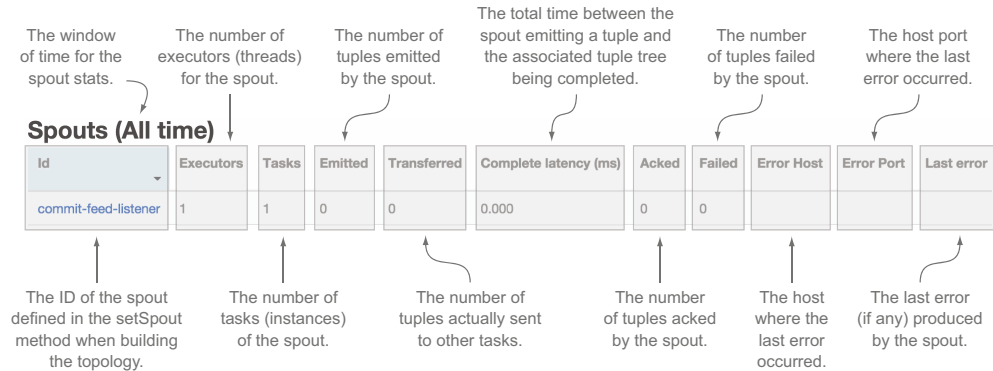


Figure 5.19 The Spouts section on the Topology summary screen of the Storm UI

BOLT STATS

The Bolts section shows the statistics for all bolts in the topology. The statistics are for the window of time that's selected in the Topology Stats section (all time, the past 10 minutes, the past 3 hours, or the past day). Figure 5.20 provides more detail up to the Capacity column. Figure 5.21 offers more detail on the remaining columns.

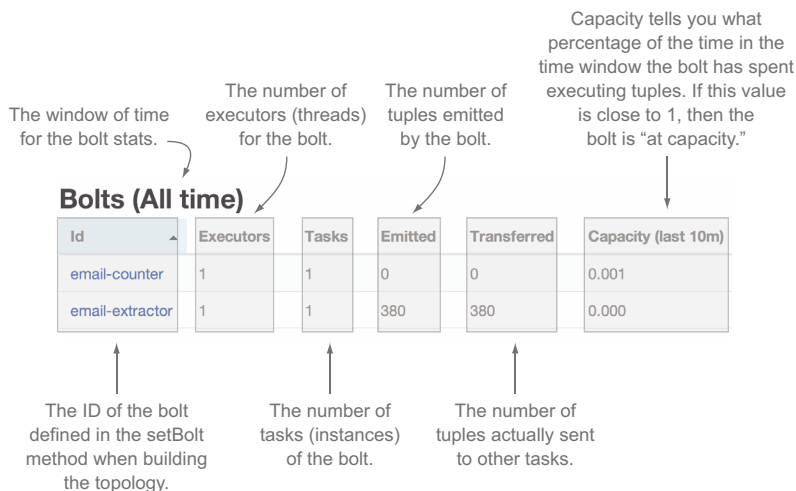


Figure 5.20 The Bolts section on the Topology summary screen of the Storm UI, up to Capacity column

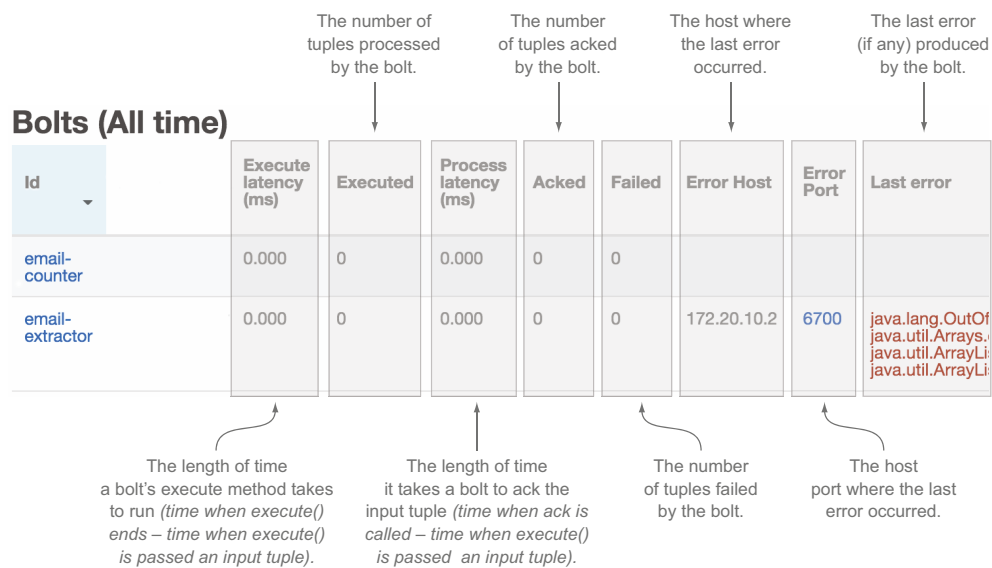


Figure 5.21 The Bolts section on the Topology summary screen of the Storm UI, remaining columns

TOPOLOGY CONFIGURATION

The Topology Configuration lists the configuration defined for the particular topology being viewed. Figure 5.22 provides more detail on what information you see in this section.

From the Topology summary screen, you can dive into one of the individual spouts or bolts. You access an individual spout or bolt by clicking on a spout or bolt name while on the Topology summary screen.

Topology Configuration

Key	Value
dev.zookeeper.path	/tmp/dev-storm-zookeeper
drpc.childopts	-Xmx768m
drpc.invocations.port	3773
drpc.port	3772
drpc.queue.size	128
drpc.request.timeout.secs	600
drpc.worker.threads	64

↑
The Topology
Configuration item.

↑
The value of the Topology
Configuration item defined in
defaults.yaml, storm.yaml,
or in code when building
the topology

Figure 5.22 The Topology Configuration section on the Topology summary screen of the Storm UI

5.5.3 Storm UI: individual spout/bolt summary

In the UI, an individual bolt contains six sections, as you can see in figure 5.23.

COMPONENT SUMMARY

The Component summary section shows some high-level information about the bolt or spout being observed. Figure 5.24 provides more details.

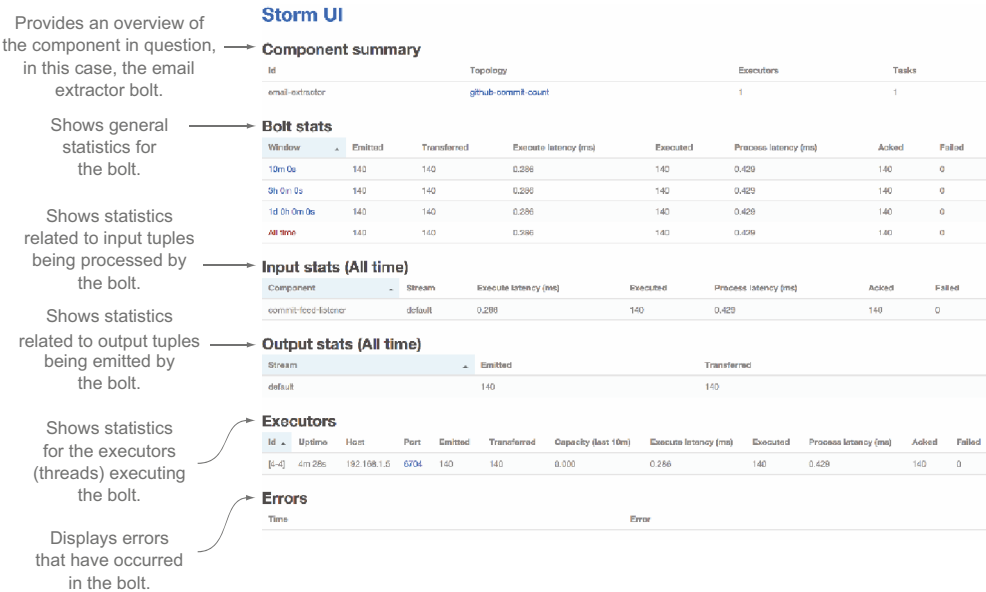


Figure 5.23 The bolt summary screen in the Storm UI

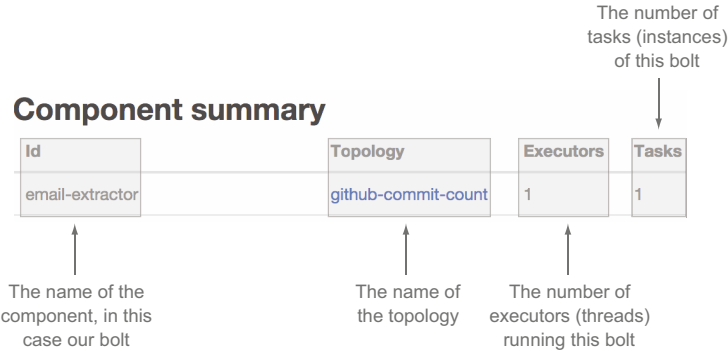


Figure 5.24 The Component Summary section for a bolt in the Storm UI

BOLT STATS

The Bolt stats section provides much of the same information that you saw in the Bolts section for the Topology summary, but the information is limited to an individual bolt (see figure 5.25).

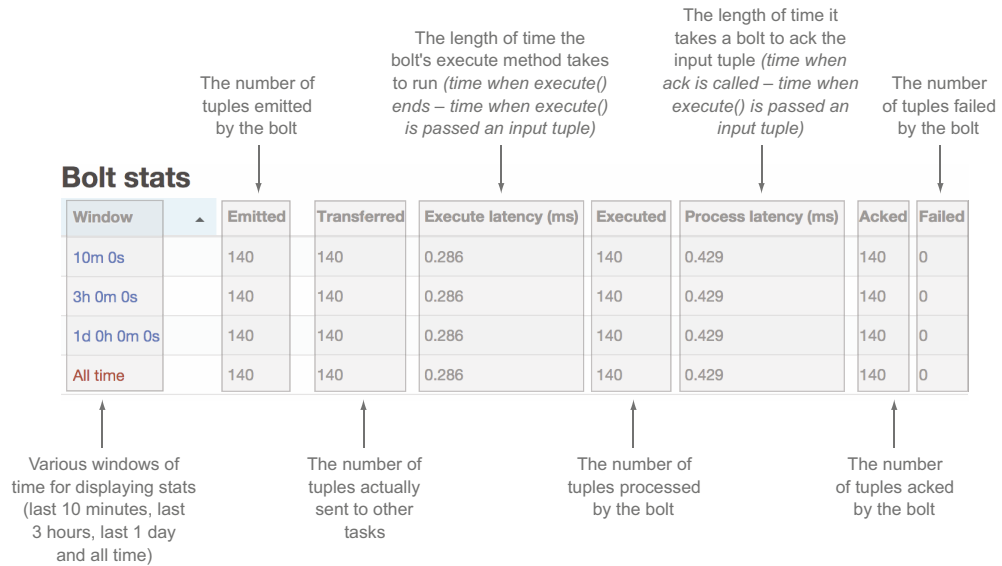


Figure 5.25 The Bolt stats section in the Storm UI

INPUT STATS

The Input stats section shows statistics related to tuples being consumed by the bolt. The statistics are relative to a particular stream; in this case it's the “default” stream. Figure 5.26 goes into more detail about this section.

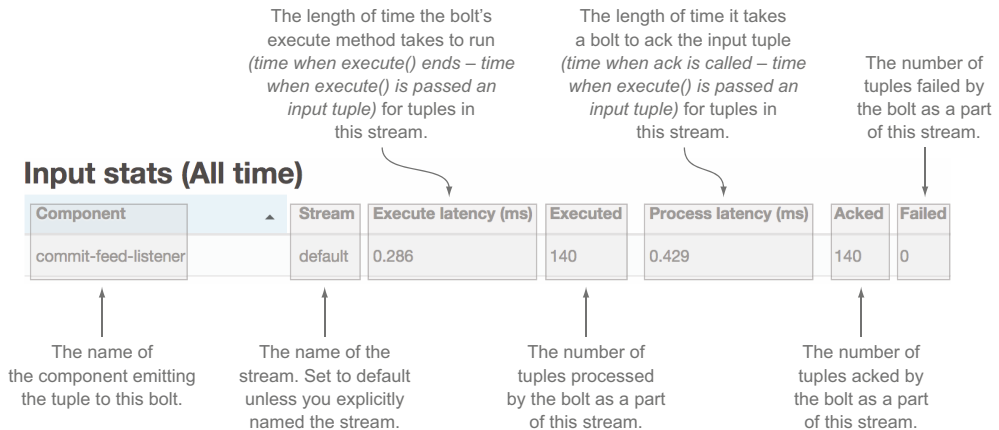


Figure 5.26 The Input stats section for a bolt in the Storm UI

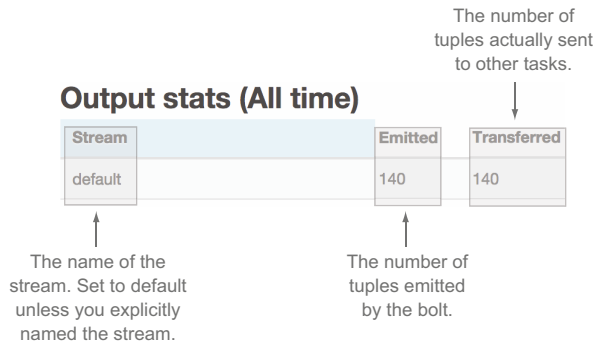


Figure 5.27 The Output stats section for a bolt in the Storm UI

OUTPUT STATS

The Output stats section shows statistics related to tuples being emitted by the bolt (see figure 5.27).

EXECUTORS

The Executors section shows the statistics for all executors running instances of the particular bolt. We’ve split this section into two figures. Figure 5.28 shows the first part and figure 5.29 shows the second.

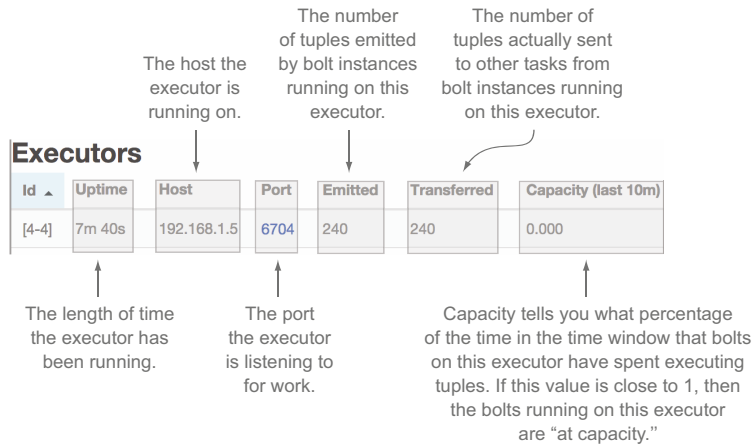


Figure 5.28 The Executors section for a bolt in the Storm UI, through Capacity column

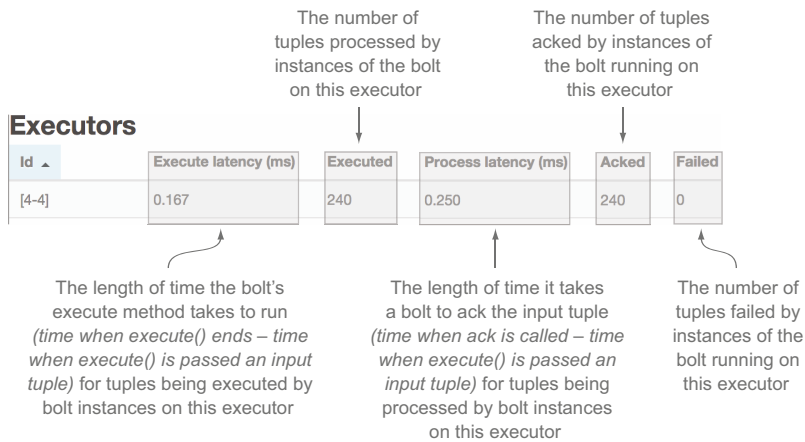


Figure 5.29 The Executors section for a bolt in the Storm UI, remaining columns

ERRORS

The Errors section shows you a history of errors experienced by this bolt, as you can see in figure 5.30.

The Storm UI provides a wealth of information, giving you a clear picture of how your topology is operating in production. With the Storm UI, you can quickly tell if your topology is healthy or if something is amiss. You should easily be able to spot errors your topology is encountering while also being able to quickly identify other issues, such as bottlenecks.

As you can probably imagine, once you deploy your topology to a production Storm cluster, your job as a developer doesn't end. Once it's deployed, you enter a whole new world of making sure your topology is running as efficiently as possible. This is the world of tuning and troubleshooting. We've devoted the next two chapters to those tasks.

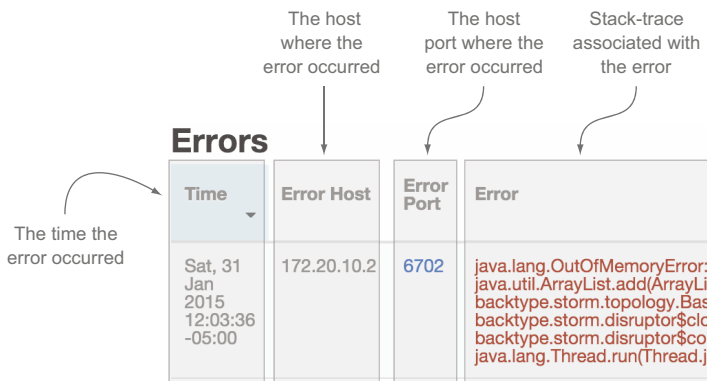


Figure 5.30 The Errors section for a bolt in the Storm UI

This chapter laid the foundation for tuning by explaining the parts of a Storm cluster and what each part of the Storm cluster does. We also provided a thorough explanation of the primary tool you'll use while tuning and troubleshooting: the Storm UI.

5.6 Summary

In this chapter you learned the following:

- A Storm cluster consists of Nimbus, which acts as the control center, and Supervisors, which execute the logic in the instances of spouts and bolts.
- A Zookeeper cluster is necessary to run alongside a Storm cluster as it coordinates communication between Nimbus/Supervisors while also maintaining state.
- Supervisors run worker processes (JVMs), which in turn run executors (threads) and tasks (instances of spouts/bolts).
- How to install a Storm cluster, including the key configuration options that must be set in order to get the cluster running.
- How to deploy your topologies to a Storm cluster and how running them on the cluster is really no different than running them locally.
- What the Storm UI is and how the different parts of the Storm ecosystem map to the different screens of the Storm UI.
- What information each section of the Storm UI provides and how these pieces of information may be useful for tuning and troubleshooting your topologies.

Storm Applied

Allen • Jankowski • Pathirana



It's hard to make sense out of data when it's coming at you fast. Like Hadoop, Storm processes large amounts of data but it does it reliably and in real time, guaranteeing that every message will be processed. Storm allows you to scale with your data as it grows, making it an excellent platform to solve your big data problems.

Storm Applied is an example-driven guide to processing and analyzing real-time data streams. This immediately useful book starts by teaching you how to design Storm solutions the right way. Then, it quickly dives into real-world case studies that show you how to scale a high-throughput stream processor, ensure smooth operation within a production cluster, and more. Along the way, you'll learn to use Trident for stateful stream processing, along with other tools from the Storm ecosystem.

What's Inside

- Mapping real problems to Storm components
- Performance tuning and scaling
- Practical troubleshooting and debugging
- Exactly-once processing with Trident

This book moves through the basics quickly. While prior experience with Storm is not assumed, some experience with big data and real-time systems is helpful.

Sean T. Allen, Matthew Jankowski, and Peter Pathirana lead the development team for a high-volume, search-intensive commercial web application at TheLadders.

To download their free eBook in PDF, ePub, and Kindle formats, owners of this book should visit manning.com/StormApplied

“Will no doubt become the definitive practitioner’s guide for Storm users.”

—From the Foreword by Andrew Montalenti

“The book’s practical approach to Storm will save you a lot of hassle and a lot of time.”

—Tanguy Leroux, Elasticsearch

“Great introduction to distributed computing with lots of real-world examples.”

—Shay Elkin, Tangent Logic

“Go beyond the MapReduce way of thinking to solve big data problems.”

—Muthusamy Manigandan OzoneMedia

